Structuring a Machine Learning Project

I. Definition

Domain
Liquor sales data in Iowa.
Of a specific subset of licensees though the state of Iowa, 2015 - Q1 2016.

Problem Statement
Predict 2016 sales by creating a regression model (e.g. OLS, SGD, Lasso, Ridge, ElasticNet) using Q1 2015 sales and the entire year to create a model for the entire year of 2016.

Predict the optimal opening location of a new store by creating a regression model (e.g. OLS, SGD, Lasso, Ridge, ElasticNet) using same sales data. Or by identifying prime location using all data.

Metric
R^2 (Coefficient of Determination), Mean Squared Error, Mean Absolute Error

II. Preparation

Data Exploration
Read the README. What types of data (ints, floats, strings, booleans, etc)/ `astype`. What is necessary/extraneous? Pivot tables. `groupby`, `agg`, check for garbage data (GIGO), measures of central tendency, densities, `describe()`, `info`, `unique`, `shape`, recategorization, sort,

Visualization
`sns.barplot`, `sns.heatmap`, `sns.pairplot`, `pd.scatterplot`, `sns.countplot`, `plt.hist`, `plt.plot`, `sns.stripplot`, `sns.swarmplot`, `sns.violinplot`

Algorithms and Techniques
Ordinary Least Squares, optionally apply regularization using l2norm (ridge), l1norm (lasso), both (elastic).
Stochastic Gradient Descent, optionally apply regularization
Pivot Tables and groupby-aggregate function, merge-join,
Heavy emphasis on data processing, large messy dataset

Benchmark
Performance with respect to metric selected above.

Code Design
1. Load the Data in Pandas
2. Clean/Munge the Data - get rid/impute nans, get rid of errant strings,
3. Verify the quality of the data - are there errors? missing data/data completeness? data mismatch? data over/under representation
4. Preprocess the Raw Data making calculated/derived fields (e.g. log transform, poly transform), dummy variables, casting to numerical types, reduce dimensionality
5. Feature Selection by Forward Selection, Backward Selection, Grid Search
6. Split the data into training, test sets
7. Normalization of the data
8. Create a new model
9. Fit the model
10. Score the model

III. Methodology

Data Preprocessing
1. Verify the quality of the data - are there errors? missing data/data completeness? data mismatch? data over/under representation
2. Preprocess the Raw Data making calculated/derived fields (e.g. log transform, poly transform), dummy variables, casting to numerical types, reduce dimensionality
3. Feature Selection by Forward Selection, Backward Selection, Grid Search

Implementation
load the data from csv using pandas
Visualize using seaborn
Model and Score using sklearn

Refinement
Tried, not so good
tried OLS which gave negative r squared
fit on overly correlated features (no new information)
added another dataset (unemployment data), added v little
Created dummy variables to include categorial data
Benchmark not clear (maybe some articles on the site)

Worked
Premium alcohol was a good feature
Identifying by type of alcohol
number of stores
pandas pivot tables worked well

IV. Results

Model Evaluation and Validation
Rather than sales, look at changes in sales over time in order to identify where the greatest demand will be. 520xx
(delta sales/zip_code/year_i) -> lasso
use built model to predict on full dataset, looked for largest delta

Justification

V. Conclusion

Reflection

Visualization

Recommendations